



**30 SEPTIEMBRE 2017**



**la térmica**  
CENTRO DE CULTURA CONTEMPORÁNEA






30 SEPTIEMBRE 2017



# DEEP LEARNING CON CNTK

## Taller de Visión Artificial



# Manuel Rodrigo Cabello Malagón

Software Engineer @ PLAIN CONCEPTS

@mrcabellom 

mrcabellom@Gmail.com 

<https://geeks.ms/mrcabellom/> 




# Pablo Doval

DATA PONTIFEX @ PLAIN CONCEPTS


*"I work with code and data,  
but don't tell my mom; she  
thinks I'm a piano player in  
a whorehouse."*

# ¿QUÉ ES ESTO?






**¿PARA QUE NOS SIRVE ESTO?**



sheep  
0.84 - > 0.96



sheep  
0.99



sheep  
0.93













# MACHINE LEARNING EN AZURE



## Azure ML

- Fácil y sencillo
- Escalabilidad limitada
- Version Renovada hace 4 días ;)



## R Server

- Lenguaje conocido
- Escalado a grandes datasets
- Evaluaciones en tiempo real



## HDInsight

- Mahout
- SparkML
- Escalabilidad casi ilimitada



## Cognitive Services

- APIs listas para usar
- Modelos pre-entrenados
- No son necesarios conocimientos de ML



## CNTK

- Modelos complejos
- Entrenamiento personalizado
- Entrenamiento multi server/multi GPU

A large group of approximately 50 employees, mostly men, are posing for a group photo in front of a modern, multi-story brick building. They are arranged in several rows, with some standing and some kneeling or sitting in the front. Most are wearing dark blue t-shirts with the 'plain concepts' logo. The building behind them has large windows and a palm tree is visible on the left. The 'plain concepts' logo is overlaid in the bottom right corner.

★★★★★

Vision

Speech

Language

Knowledge

Search



I am not really confident, but I think it's a group of people on a sidewalk and they seem



How did I do?



## Vision

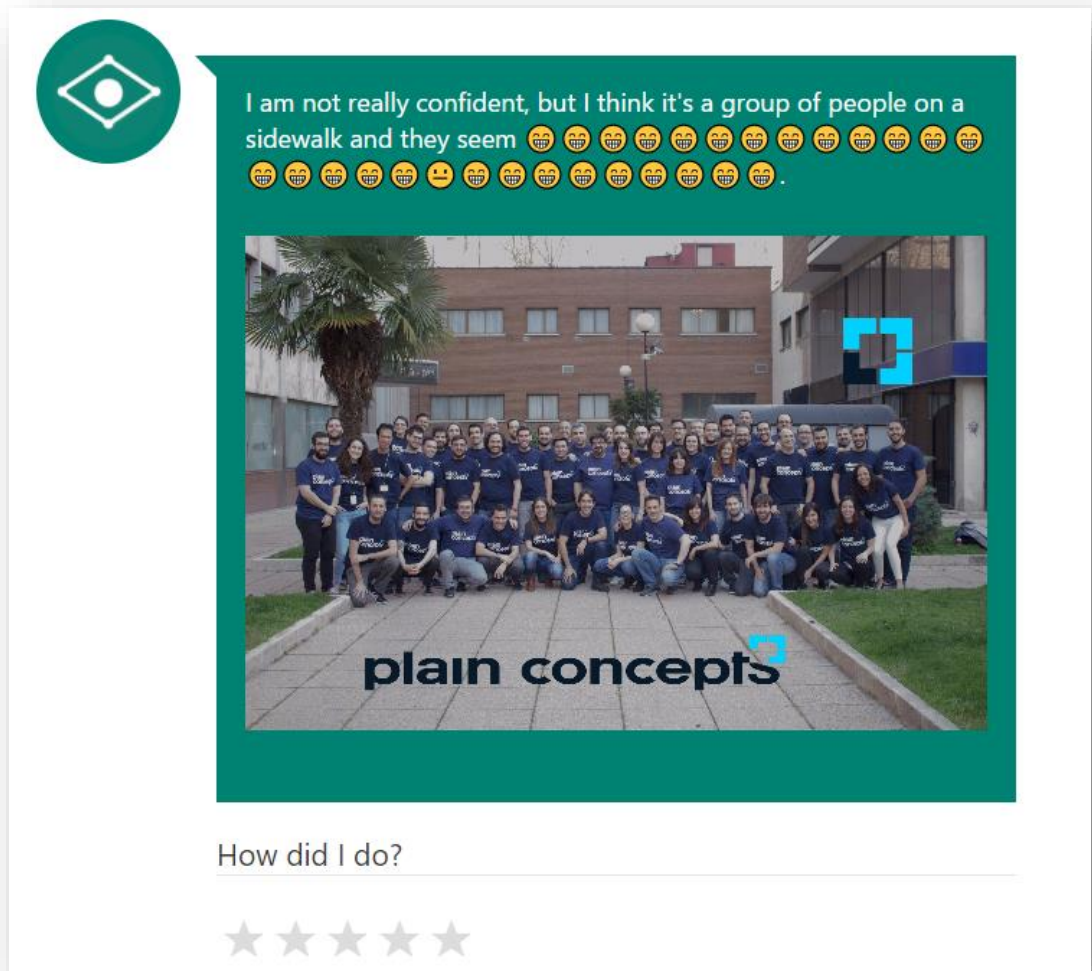
- Computer Vision
- Content Moderator
- Emotion API
- Face API
- ...

## Speech

## Language

## Knowledge

Search



I am not really confident, but I think it's a group of people on a sidewalk and they seem



How did I do?





## Vision

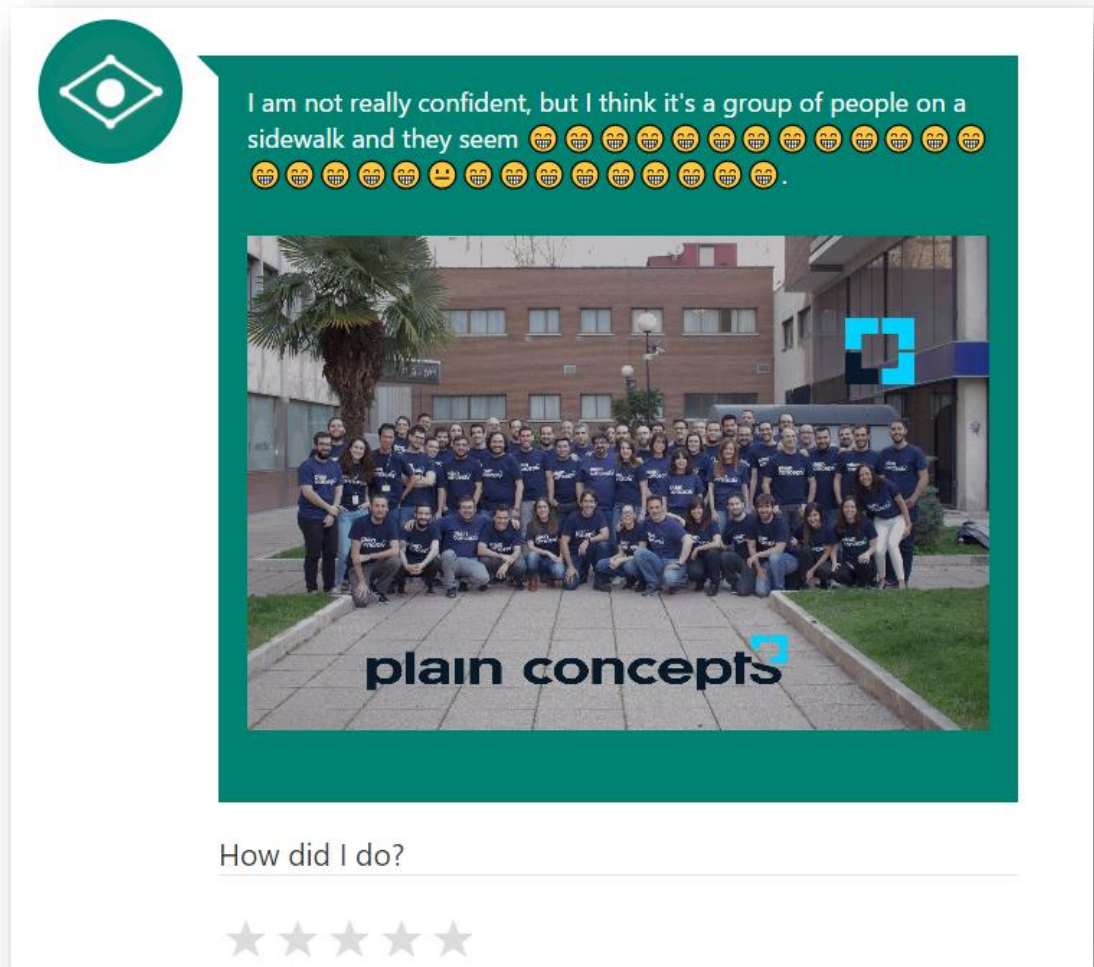
## Speech

- Bing Speech API
- Custom Speech Service
- Speaker Recognition Service

## Language

Knowledge

Search



I am not really confident, but I think it's a group of people on a sidewalk and they seem .



How did I do?

## Vision

## Speech

## Language

- Spell Check API
- LUIS
- Linguistic Analysis
- Translator API
- ...

## Knowledge

Search



I am not really confident, but I think it's a group of people on a sidewalk and they seem   
.



How did I do?



## Vision

## Speech

## Language

## Knowledge

- Entity Linking Service
- Academic Knowledge Service
- Q&A Maker
- ...

Search



I am not really confident, but I think it's a group of people on a sidewalk and they seem .



How did I do?



# Vision

# Speech

# Language

# Knowledge

# Search

- Autosuggest API
- Image Search API
- News Search API
- ...

[illegible]

How did I do?



# ¿CUANDO **NO** USAR COGNITIVE SERVICES?





# ¿CUANDO **NO** USAR COGNITIVE SERVICES?





# ¿CUANDO **NO** USAR COGNITIVE SERVICES?



# ¿CUANDO **NO** USAR COGNITIVE SERVICES?





# **DEEP LEARNING CON CNTK**



# **DEEP LEARNING** CON CNTK

# DEEP LEARNING

---

Machine Learning

---

Basados en Redes Neuronales

---

Más de una capa oculta (*Deep Models*)

---

A large, light gray gear graphic is positioned in the top-left corner of the slide, partially cut off by the edge.

# DEEP LEARNING CON **CNTK**



# CNTK

¿Por qué?

- Desarrollo más rápido de los modelos
- Uso de múltiples CPUs y GPUs
- Multi-plataforma
- Es el producto que Microsoft usa internamente

# CNTK

- Es el acrónimo del algo.
- Es un framework de Deep Learning, y es open-source, con contribuciones externas relevantes del MIT y Stanford, entre otros.
- Permite expresar redes neuronales complejas de un modo relativamente sencillo, y se encarga de todos los pasos de su ciclo de vida: desde el entrenamiento hasta la evaluación.
- Desarrollo en BrainScript, Python (.NET en camino)



# **LAB 0: INSTALACIÓN**




# **TÉCNICAS DE VISIÓN ARTIFICIAL**

## **Regresión Logística**


# UN EJEMPLO

0	0	0	0	0	0	0	0	0
0	1	0	0	0	0	0	1	0
0	0	1	0	0	0	1	0	0
0	0	0	1	0	1	0	0	0
0	0	0	0	1	0	0	0	0
0	0	0	1	0	1	0	0	0
0	0	1	0	0	0	1	0	0
0	1	0	0	0	0	0	1	0
0	0	0	0	0	0	0	0	0





"X"

0	0	0	0	0	0	0	0	0
0	0	0	1	1	1	0	0	0
0	0	1	0	0	0	1	0	0
0	1	0	0	0	0	0	1	0
0	1	0	0	0	0	0	1	0
0	1	0	0	0	0	0	1	0
0	0	1	0	0	0	1	0	0
0	0	0	1	1	1	0	0	0
0	0	0	0	0	0	0	0	0



"O"



Bias ( $\vec{b}$ )






$\Sigma$  Sum (weights x pixels) =  $\vec{w}_x \cdot \vec{x}$

$\Sigma$  Sum (weights x pixels) =  $\vec{w}_o \cdot \vec{x}$



$$z = \vec{w} \cdot \vec{x} + b$$



# SIGMOIDE




# SOFTMAX




$$s(y_i) = \frac{e^{y_i}}{\sum_j e^{y_j}}$$

Esto ya son probabilidades


# FUNCIÓN DE PERDIDA



One-hot encoding



# FUNCIÓN DE PERDIDA




1 0

$$\text{error} = \frac{1}{n} \sum_{i=1}^n (y_j - p_j)^2$$

0.75 0.25

# CROSS ENTROPY




1 0

$$\text{error} = y_j^T \log P_j$$

0.75 0.25

# ENTRENAMIENTO

$x$  = vector de pixels de entrada




$y$  = vector de etiquetas  
(one-hot)

(1, 0)

(0, 1)


(1, 0)

```
z = times(x, W) + b
cross_entropy_with_softmax(z, y)
classification_error(z, y)
```




Training Data

# DESCENSO DE GRADIENTE




# DESCENSO DE GRADIENTE





# DESCENSO DE GRADIENTE






# **LAB 1: REGRESIÓN LOGÍSTICA**




# **TÉCNICAS DE VISIÓN ARTIFICIAL**

## **Perceptrón Multicapa**


# REDES NEURONALES - CLASIFICADOR



# REDES NEURONALES - CLASIFICADOR




# FORWARD PROPAGATION







# FORWARD PROPAGATION




# FORWARD PROPAGATION



# FORWARD PROPAGATION




# PESOS Y BIAS




# ENTRENAMIENTO

$$\text{Coste} = \text{Valor Generado} - \text{Valor Real}$$



# BACK PROPAGATION

$$\text{Coste} = \text{Valor Generado} - \text{Valor Real}$$








# **LAB 2: PERCEPTRÓN MULTICAPA**



# DEEP LEARNING

# IMAGENET CHALLENGE





# **TÉCNICAS DE VISIÓN ARTIFICIAL**


## **Redes Neuronales Convolucionales**



# **TÉCNICAS DE VISIÓN ARTIFICIAL**




## **Redes Neuronales Convolucionales**

### **Con Excel**










# FILTRADO



# FILTRADO



Filtro 1   Filtro 2   Filtro 3



1.0

0.11

0.55

# FILTRADO

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

1	-1	-1
-1	1	-1
-1	-1	1

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

# CONVOLUCIÓN

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1



1	-1	-1
-1	1	-1
-1	-1	1

=

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

# CONVOLUCIÓN

-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1
-1	-1	1	-1	-1	-1	1	-1
-1	-1	-1	1	-1	1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1
-1	-1	1	-1	-1	-1	1	-1
-1	1	-1	-1	-1	-1	-1	1
-1	-1	-1	-1	-1	-1	-1	-1



1	-1	-1
-1	1	-1
-1	-1	1

=

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1
-1	-1	1	-1	-1	-1	1	-1
-1	-1	-1	1	-1	1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1
-1	-1	1	-1	-1	-1	1	-1
-1	1	-1	-1	-1	-1	-1	1
-1	-1	-1	-1	-1	-1	-1	-1



-1	-1	1
-1	1	-1
1	-1	-1

=

0.33	-0.11	0.55	0.33	0.11	-0.11	0.77
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.77	-0.11	0.11	0.33	0.55	-0.11	0.33

# CONVOLUCIÓN

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1



1	-1	-1
-1	1	-1
-1	-1	1

1	-1	1
-1	1	-1
1	-1	1

-1	-1	1
-1	1	-1
1	-1	-1

=

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

0.33	-0.55	-0.11	-0.11	0.11	-0.55	0.33
-0.55	0.55	-0.55	0.33	-0.55	0.55	-0.55
0.11	-0.55	0.55	-0.77	0.55	-0.55	0.11
-0.11	0.33	-0.77	1.00	-0.77	0.33	-0.11
0.11	-0.55	0.55	-0.77	0.55	-0.55	0.11
-0.55	0.55	-0.55	0.33	-0.55	0.55	-0.55
0.33	-0.55	0.11	-0.11	0.11	-0.55	0.33

0.33	-0.11	0.55	0.33	0.11	-0.11	0.77
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.77	-0.11	0.11	0.33	0.55	-0.11	0.33

# POOLING

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

# POOLING

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77



# AVERAGE POOLING – STRIDE 1

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

0.39	0.22	0.17	0.28	0.11	0.06
0.22	0.45	0.22	0.00	0.00	0.11
0.17	0.22	0.22	0.00	0.00	0.28
0.28	0.00	0.00	0.22	0.22	0.17
0.11	0.00	0.00	0.22	0.45	0.22
0.06	0.11	0.28	0.17	0.22	0.39

# MAX POOLING – STRIDE 1

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

1.00	1.00	0.33	0.55	0.55	0.33
1.00	1.00	1.00	0.33	0.11	0.55
0.33	1.00	1.00	0.55	0.33	0.55
0.55	0.33	0.55	1.00	1.00	0.33
0.55	0.11	0.33	1.00	1.00	1.00
0.33	0.55	0.55	0.33	1.00	1.00


# MAX POOLING – STRIDE 2

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77


1.00	0.33	0.55	0.33
0.33	1.00	0.33	0.55
0.55	0.33	1.00	0.11
0.33	0.55	0.11	0.77

# POOLING - ¿PARA QUE?

# FUNCIONES DE ACTIVACIÓN




# FUNCIONES DE ACTIVACIÓN



## FUNCIÓN LOGÍSTICA (“*sigmoide*”):

- Simple.
- Pegas:
  - Saturación
  - No centrada en cero
- No la usamos, pero le agradecemos los servicios prestados :)


# FUNCIONES DE ACTIVACIÓN



## TANGENTE HIPERBÓLICA ( $\tanh$ ):

- Resuelve el problema del no-centrado en cero
- Sigue 'matando gradientes' por saturación
- Siempre es preferible a la función logística

# FUNCIONES DE ACTIVACIÓN




## RECTIFIED LINEAR UNITS (ReLU):

- Optimización para la convergencia del SGD
- Computacionalmente más sencilla (no hay que exponenciar, etc..)
- Frágiles (pueden 'morir' dejando parte de la red inaccesible)



# FUNCIONES DE ACTIVACIÓN



## LEAKY ReLU:

- Más estables que ReLU tradicional
- La pendiente de la región negativa puede parametrizarse
- Caso particular del *MaxOut*



# ReLU

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77





0.77	0.00	0.11	0.33	0.55	0.00	0.33
0.00	1.00	0.00	0.33	0.00	0.11	0.00
0.11	0.00	1.00	0.00	0.11	0.00	0.55
0.33	0.33	0.00	0.55	0.00	0.33	0.33
0.55	0.00	0.11	0.00	1.00	0.00	0.11
0.00	0.11	0.00	0.33	0.00	1.00	0.00
0.33	0.00	0.55	0.33	0.11	0.00	0.77

# DROPOUT



# DROPOUT



# ¿QUÉ FALTA AQUÍ?

Filtro 1      Filtro 2      Filtro 3


1	-1	-1
-1	1	-1
-1	-1	1

-1	-1	1
-1	1	-1
1	-1	-1


1	-1	1
-1	1	-1
1	-1	1

¿De dónde han salido estos?



Faces






Cars



Elephants



Chairs




# **LAB 2: REDES NEURONALES CONVOLUCIONALES**


# ALEXNET EN CNTK (PYTHON)



# EJEMPLO: AlexNet



# EJEMPLO: AlexNet



# PRODUCCIONALIZACIÓN

- Nuevo Servicio de AzureML
- Maquinas Virtuales de Azure de tipo NC
  - Coordinación con Azure Batch / Azure AI Training Service
- Clústers de HDInsight mediante PySpark
- Checkpoints

# ALGUNAS IDEAS

- AutoEncoders
- Recurrent Neural Networks (RNN)
- Long Short Term Memory (LSTM)

# ALGUNOS PAPERS INTERESANTES

"ImageNet Classification with Deep Convolutional Neural Networks"

[Krizhevsky, Sutskever, Hinton]

arXiv:

"Visualizing and Understanding Convolutional Networks":

[Zeiler, Fergus]

arXiv: 1311.2901v3

"Optimization Methods for Large-Scale Machine Learning":

[Bottou, Curtis, Nocedal]

arXiv: 1606.04838v2

# CALL TO ACTION

## ¡Cacharread!

Cognitive toolkit @ <https://cntk.ai>

Tutoriales @ <https://cntk.ai/pythondocs/tutorials.html>

Notebooks @ <https://notebooks.azure.com/cntk/libraries/tutorials>

"Road to Deep Learning" en Plain<TV>

Entended el backpropagation, ¡en serio! 😊